

Multiple regression

David Puelz

The University of Austin

One predictor is not enough

Holding something fixed

Is the model any good?

Do rich places vote Republican?

Do rich places vote Republican?

Everyone is sure they know the answer.
Almost everyone is wrong.



- Every U.S. county that reported in 2024: 3,107 of them
- **trump_share** – percent of the county's presidential vote for Trump
- **median_income** – median household income (\$000)
- **pct_college** – percent of adults with a bachelor's degree or higher

Election returns from county-level 2024 returns; county characteristics from the Census ACS via OpenIntro.



$$\widehat{\text{trump}\%} = 83.7 - 0.314 (\text{income in \$000})$$

Every extra \$10,000 of median income predicts about **3.1 points less** for Trump.



$$\widehat{\text{trump}\%} = 83.7 - 0.314 (\text{income in \$000})$$

Every extra \$10,000 of median income predicts about **3.1 points less** for Trump.

A tidy story. Money votes blue.

But richer counties are also better educated



Predictor	Slope	R^2
median income	-0.314 points per \$1,000	0.079
% college	-0.933 points per point	0.32

They are correlated: $r = 0.71$.



Predictor	Slope	R^2
median income	-0.314 points per \$1,000	0.079
% college	-0.933 points per point	0.32

They are correlated: $r = 0.71$.

So which one is actually doing the work?

One predictor is not enough

Holding something fixed

Is the model any good?



One predictor:

$$Y_i = \beta_0 + \beta_1(\text{income}) + \epsilon_i$$

Two predictors:

$$Y_i = \beta_0 + \beta_1(\text{income}) + \beta_2(\% \text{ college}) + \epsilon_i$$



One predictor:

$$Y_i = \beta_0 + \beta_1(\text{income}) + \epsilon_i$$

Two predictors:

$$Y_i = \beta_0 + \beta_1(\text{income}) + \beta_2(\% \text{ college}) + \epsilon_i$$

Same idea. Still least squares. We just choose three numbers instead of two.

```
model <- lm(trump ~ inc + ba, data=counties)
summary(model)
```

```
Call:
lm(formula = trump ~ inc + ba, data = counties)
```

```
Residuals:
```

	Min	1Q	Median	3Q	Max
	-61.967	-6.841	1.452	8.528	44.910

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	79.29733	0.90047	88.06	<2e-16 ***
inc	0.26630	0.02289	11.63	<2e-16 ***
ba	-1.21066	0.03376	-35.86	<2e-16 ***

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 12.71 on 3104 degrees of freedom
```

```
Multiple R-squared:  0.3489, Adjusted R-squared:  0.3485
```

```
F-statistic: 831.6 on 2 and 3104 DF,  p-value: < 2.2e-16
```



$$\widehat{\text{trump\%}} = 79.3 + 0.266 (\text{income}) - 1.211 (\% \text{ college})$$

Alone, income was worth -0.314 points per \$1,000.

Next to education, it is worth $+0.266$.




$$\widehat{\text{trump}\%} = 79.3 + 0.266 (\text{income}) - 1.211 (\% \text{ college})$$


Alone, income was worth -0.314 points per \$1,000.

Next to education, it is worth $+0.266$.


Same counties. Same election. What changed?




Only counties that are 10–15% college




Only counties that are 18–23% college



Only counties that are 26–31% college







- The **dashed line** is the simple regression across all 3,107 counties. It slopes **down**.
- The **colored lines** hold education roughly fixed. They slope **up** (or flat).
- Compare two counties at different incomes in the raw data and you are also comparing a less-educated place to a more-educated one. You slide across the whole picture, not along one line.
- Hold education fixed and the sign reverses: $\hat{\beta}_1 = +0.266$. **Among similarly educated counties, richer ones voted more for Trump.**



- Holding the other fixed, income pushes Trump share **up** and education pushes it **down**. But richer counties are also more educated ($r = 0.71$), so leaving education out lets income soak up education's negative effect.
- Both slopes are correct. They answer **different questions**: “richer places” versus “richer places with the same education.”
- **A caution**: these are counties, not people. A pattern across counties need not hold for individual voters. That mistake has a name, the **ecological fallacy**, and it is everywhere in political commentary.

One predictor is not enough

Holding something fixed

Is the model any good?



Predictors	R^2	s_e
income	0.079	15.12
% college	0.32	12.99
both	0.349	12.71

- R^2 **always rises** when you add a variable, and the values **do not add**: $0.079 + 0.32 = 0.4$, but together they give only 0.349. The predictors overlap.
- $\sqrt{R^2} = 0.591$ is exactly the correlation between $\hat{Y}$ and Y .
- $s_e = 12.71$ points is the typical miss on one county. R also prints **adjusted** R^2 , which penalizes extra parameters.



- **Overall null:** $\beta_1 = \beta_2 = 0$, equivalently $R^2 = 0$ in the population, equivalently the model beats guessing $\bar{Y}$ by nothing. The F -test gives $F = 832$, $p < 2.2 \times 10^{-16}$. Rejected.
- **Individual coefficients:** test $H_0 : \beta_i = 0$ one at a time. This asks whether a predictor adds anything **once the other is in**. Income: $p = 1.2 \times 10^{-30}$. Education: $p = 6.6 \times 10^{-236}$.
- A smaller p -value does **not** mean a stronger predictor.



- Adding income to the education model is highly significant, yet R^2 moves only from 0.32 to 0.349.
- As a **prediction** tool it earns very little.
- But it completely reversed what we would say about money and voting. Sometimes a variable belongs in the model not to predict, but **to hold something fixed**.
- **There is no universal rule for adding or removing variables.**

A slope is never just a number.

It is an answer to the question
“compared to what?”